

# Introduction to parallel computing with R

Hana Ševčíková

University of Washington  
hanas@uw.edu

useR!2017 Conference, Brussels, July 4–7, 2017

# Motivation

Goal:

# Motivation

**Goal:** Solve a given problem faster by utilizing more resources.

# Motivation

**Goal:** Solve a given problem faster by utilizing more resources.



# Motivation

**Goal:** Solve a given problem faster by utilizing more resources.



# Motivation

**Goal:** Solve a given problem faster by utilizing more resources.



Serial processing

# Motivation

**Goal:** Solve a given problem faster by utilizing more resources.



Serial processing



Parallel processing

# Hardware Considerations



# Hardware Considerations

- ▶ Central processing unit (CPU):

# Hardware Considerations

- ▶ Central processing unit (CPU):
  - ▶ Multi-processor (CPU, core) computer

# Hardware Considerations

- ▶ Central processing unit (CPU):
  - ▶ Multi-processor (CPU, core) computer
  - ▶ Cluster of single- or multi-processors computers


# Hardware Considerations

- ▶ Central processing unit (CPU):
  - ▶ Multi-processor (CPU, core) computer
  - ▶ Cluster of single- or multi-processors computers
- ▶ Memory:

# Hardware Considerations

- ▶ Central processing unit (CPU):
  - ▶ Multi-processor (CPU, core) computer
  - ▶ Cluster of single- or multi-processors computers
- ▶ Memory:

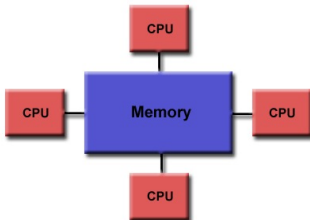
Shared Memory



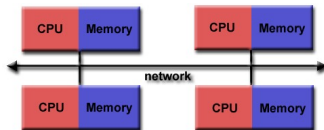
# Hardware Considerations

- ▶ Central processing unit (CPU):
  - ▶ Multi-processor (CPU, core) computer
  - ▶ Cluster of single- or multi-processors computers
- ▶ Memory:

Shared Memory



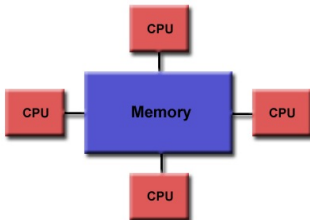
Distributed Memory




# Hardware Considerations

- ▶ Central processing unit (CPU):
  - ▶ Multi-processor (CPU, core) computer
  - ▶ Cluster of single- or multi-processors computers
- ▶ Memory:


Shared Memory



Distributed Memory



Hybrid Memory



# Parallel Programming

- ▶ Programming Paradigms:



# Parallel Programming

- ▶ Programming Paradigms:
  - ▶ Implicit Parallelism:

# Parallel Programming

- ▶ Programming Paradigms:
  - ▶ Implicit Parallelism:
    - ▶ Does not need directives for parallelization from user

# Parallel Programming

- ▶ Programming Paradigms:
  - ▶ Implicit Parallelism:
    - ▶ Does not need directives for parallelization from user
    - ▶ Examples: parallel compilers, HPF, OpenMP, ScaLAPACK, PBLAS

# Parallel Programming

- ▶ **Programming Paradigms:**
  - ▶ **Implicit Parallelism:**
    - ▶ Does not need directives for parallelization from user
    - ▶ Examples: parallel compilers, HPF, OpenMP, ScaLAPACK, PBLAS
  - ▶ **Explicit Parallelism:**

# Parallel Programming

- ▶ **Programming Paradigms:**
  - ▶ **Implicit Parallelism:**
    - ▶ Does not need directives for parallelization from user
    - ▶ Examples: parallel compilers, HPF, OpenMP, ScaLAPACK, PBLAS
  - ▶ **Explicit Parallelism:**
    - ▶ Message-passing libraries (MPI)

# Parallel Programming

- ▶ **Programming Paradigms:**

- ▶ **Implicit Parallelism:**

- ▶ Does not need directives for parallelization from user
    - ▶ Examples: parallel compilers, HPF, OpenMP, ScaLAPACK, PBLAS

- ▶ **Explicit Parallelism:**

- ▶ Message-passing libraries (MPI)
    - ▶ Supported by various programming languages including R

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ Shared Memory Model:



# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ Shared Memory Model:
  - ▶ Processes share a common address space which they can read & write to.

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**



# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**
  - ▶ Master assigns jobs to a pool of workers.

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**
  - ▶ Master assigns jobs to a pool of workers. Workers perform their task independently.

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**
  - ▶ Master assigns jobs to a pool of workers. Workers perform their task independently. Master collects results.

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**
  - ▶ Master assigns jobs to a pool of workers. Workers perform their task independently. Master collects results.
  - ▶ Suited for *embarrassingly parallel* applications.

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**
  - ▶ Master assigns jobs to a pool of workers. Workers perform their task independently. Master collects results.
  - ▶ Suited for *embarrassingly parallel* applications.
  - ▶ R: snow (now partly in parallel), snowFT, snowfall, foreach

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**
  - ▶ Master assigns jobs to a pool of workers. Workers perform their task independently. Master collects results.
  - ▶ Suited for *embarrassingly parallel* applications.
  - ▶ R: snow (now partly in parallel), snowFT, snowfall, foreach
- ▶ **Single Program Multiple Data (SPMD) :**

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**
  - ▶ Master assigns jobs to a pool of workers. Workers perform their task independently. Master collects results.
  - ▶ Suited for *embarrassingly parallel* applications.
  - ▶ R: snow (now partly in parallel), snowFT, snowfall, foreach
- ▶ **Single Program Multiple Data (SPMD) :**
  - ▶ Emphasis on data parallelism.

# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**
  - ▶ Master assigns jobs to a pool of workers. Workers perform their task independently. Master collects results.
  - ▶ Suited for *embarrassingly parallel* applications.
  - ▶ R: snow (now partly in parallel), snowFT, snowfall, foreach
- ▶ **Single Program Multiple Data (SPMD) :**
  - ▶ Emphasis on data parallelism.
  - ▶ One program, processors are autonomous, no manager.



# Parallel Programming Models

- ▶ Abstraction above hardware and memory architectures.
- ▶ **Shared Memory Model:**
  - ▶ Processes share a common address space which they can read & write to.
  - ▶ Examples: POSIX Threads (Unix/Linux), OpenMP (multi-platform)
  - ▶ R: Rdsm, multicore (now in parallel), foreach, Rcpp, rJava and other packages supporting OpenMP.
- ▶ **Message Passing Model:**
  - ▶ Processes use their own memory. Communication via sending and receiving messages.
  - ▶ Examples: MPI (de-facto standard)
  - ▶ R: Rmpi, pbdMPI, pbdR
- ▶ **Master-Worker Model:**
  - ▶ Master assigns jobs to a pool of workers. Workers perform their task independently. Master collects results.
  - ▶ Suited for *embarrassingly parallel* applications.
  - ▶ R: snow (now partly in parallel), snowFT, snowfall, foreach
- ▶ **Single Program Multiple Data (SPMD) :**
  - ▶ Emphasis on data parallelism.
  - ▶ One program, processors are autonomous, no manager.
  - ▶ R: pbdR et al, RHadoop, ddR

# Structure of Statistical Simulations

Many statistical simulations have the following structure:

# Structure of Statistical Simulations

Many statistical simulations have the following structure:

```
initialize.rng(...)  
for (iteration in 1:N) {  
  result[iteration] <- myfunc(...)  
}  
process(result,...)
```

# Structure of Statistical Simulations

Many statistical simulations have the following structure:

```
initialize.rng(...)  
for (iteration in 1:N) {  
  result[iteration] <- myfunc(...)  
}  
process(result,...)
```

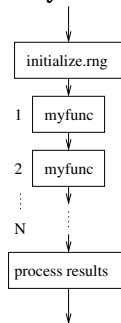
If myfunc can be processed independently, then:

# Structure of Statistical Simulations

Many statistical simulations have the following structure:

```
initialize.rng(...)  
for (iteration in 1:N) {  
  result[iteration] <- myfunc(...)  
}  
process(result,...)
```

If `myfunc` can be processed independently, then:

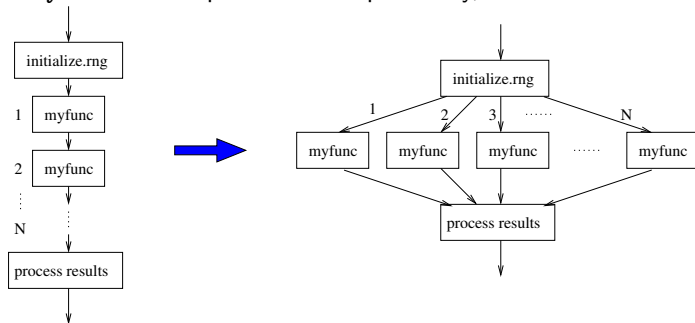


# Structure of Statistical Simulations

Many statistical simulations have the following structure:

```
initialize.rng(...)
for (iteration in 1:N) {
  result[iteration] <- myfunc(...)
}
process(result,...)
```

If `myfunc` can be processed independently, then:



# Master-Worker Paradigm

# Master-Worker Paradigm

- ▶ A cluster consists of a master and a set of workers



# Master-Worker Paradigm

- ▶ A cluster consists of a master and a set of workers (often one worker per physical node).

# Master-Worker Paradigm

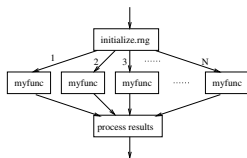
- ▶ A cluster consists of a master and a set of workers (often one worker per physical node).
- ▶ Each worker is responsible for one or more calls of `myfunc`.

# Master-Worker Paradigm

- ▶ A cluster consists of a master and a set of workers (often one worker per physical node).
- ▶ Each worker is responsible for one or more calls of `myfunc`.
- ▶ Master is responsible for sending workers everything they need and for collecting results.

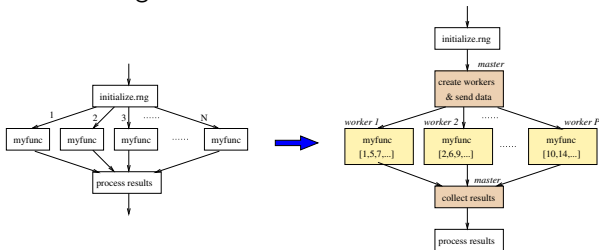
# Master-Worker Paradigm

- ▶ A cluster consists of a master and a set of workers (often one worker per physical node).
- ▶ Each worker is responsible for one or more calls of `myfunc`.
- ▶ Master is responsible for sending workers everything they need and for collecting results.



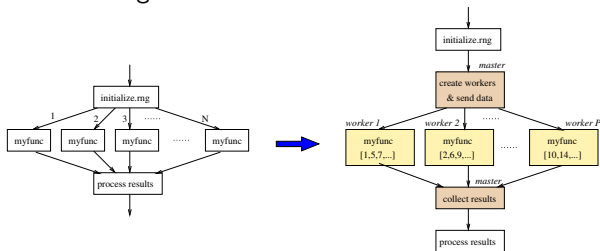
# Master-Worker Paradigm

- ▶ A cluster consists of a master and a set of workers (often one worker per physical node).
- ▶ Each worker is responsible for one or more calls of `myfunc`.
- ▶ Master is responsible for sending workers everything they need and for collecting results.



# Master-Worker Paradigm

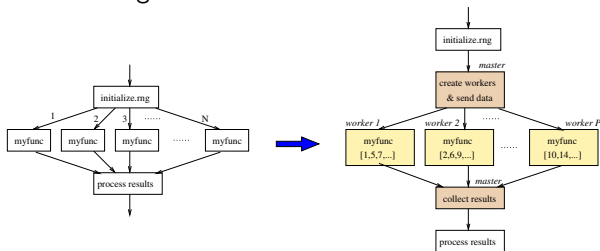
- ▶ A cluster consists of a master and a set of workers (often one worker per physical node).
- ▶ Each worker is responsible for one or more calls of `myfunc`.
- ▶ Master is responsible for sending workers everything they need and for collecting results.



- ▶ Often  $P \ll N$

# Master-Worker Paradigm

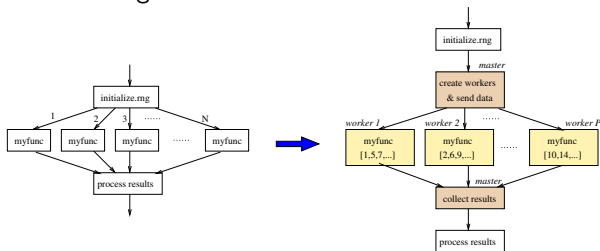
- ▶ A cluster consists of a master and a set of workers (often one worker per physical node).
- ▶ Each worker is responsible for one or more calls of `myfunc`.
- ▶ Master is responsible for sending workers everything they need and for collecting results.



- ▶ Often  $P \ll N$
- ▶ Many packages in R work in this fashion.

# Master-Worker Paradigm

- ▶ A cluster consists of a master and a set of workers (often one worker per physical node).
- ▶ Each worker is responsible for one or more calls of `myfunc`.
- ▶ Master is responsible for sending workers everything they need and for collecting results.



- ▶ Often  $P \ll N$
- ▶ Many packages in R work in this fashion.
- ▶ One of the pioneers, [snow](#) (Simple Network of Workstations) recently re-implemented as [parallel](#) (in R core).



# Random Number Generators (RNG)

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).
- ▶ Important parameters of an RNG:

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).
- ▶ Important parameters of an RNG:
  - ▶ Long period (preferably  $> 2^{100}$ )

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).
- ▶ **Important parameters** of an RNG:
  - ▶ Long period (preferably  $> 2^{100}$ )
  - ▶ Good structural (distributional) properties in high dimensions

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).
- ▶ **Important parameters** of an RNG:
  - ▶ Long period (preferably  $> 2^{100}$ )
  - ▶ Good structural (distributional) properties in high dimensions
- ▶ These parameters should hold when used in distributed environment.

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).
- ▶ **Important parameters** of an RNG:
  - ▶ Long period (preferably  $> 2^{100}$ )
  - ▶ Good structural (distributional) properties in high dimensions
- ▶ These parameters should hold when used in distributed environment.
- ▶ A good quality **RNG with multiple independent streams** proposed by L'Ecuyer et al. (2002):



# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).
- ▶ **Important parameters** of an RNG:
  - ▶ Long period (preferably  $> 2^{100}$ )
  - ▶ Good structural (distributional) properties in high dimensions
- ▶ These parameters should hold when used in distributed environment.
- ▶ A good quality **RNG with multiple independent streams** proposed by L'Ecuyer et al. (2002):
  - ▶ Period  $2^{191}$ ; streams have seeds  $2^{127}$  steps apart.

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).
- ▶ **Important parameters** of an RNG:
  - ▶ Long period (preferably  $> 2^{100}$ )
  - ▶ Good structural (distributional) properties in high dimensions
- ▶ These parameters should hold when used in distributed environment.
- ▶ A good quality **RNG with multiple independent streams** proposed by L'Ecuyer et al. (2002):
  - ▶ Period  $2^{191}$ ; streams have seeds  $2^{127}$  steps apart.
  - ▶ Parallel parts of users computation can run independent and reproducible streams.

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).
- ▶ **Important parameters** of an RNG:
  - ▶ Long period (preferably  $> 2^{100}$ )
  - ▶ Good structural (distributional) properties in high dimensions
- ▶ These parameters should hold when used in distributed environment.
- ▶ A good quality **RNG with multiple independent streams** proposed by L'Ecuyer et al. (2002):
  - ▶ Period  $2^{191}$ ; streams have seeds  $2^{127}$  steps apart.
  - ▶ Parallel parts of users computation can run independent and reproducible streams.
  - ▶ Direct interface in R: **rlecuyer**, **rstream**

# Random Number Generators (RNG)

- ▶ Many statistical applications involve random numbers (e.g. MCMCs in Bayesian methods, bootstrap, simulations).
- ▶ **Important parameters** of an RNG:
  - ▶ Long period (preferably  $> 2^{100}$ )
  - ▶ Good structural (distributional) properties in high dimensions
- ▶ These parameters should hold when used in distributed environment.
- ▶ A good quality **RNG with multiple independent streams** proposed by L'Ecuyer et al. (2002):
  - ▶ Period  $2^{191}$ ; streams have seeds  $2^{127}$  steps apart.
  - ▶ Parallel parts of users computation can run independent and reproducible streams.
  - ▶ Direct interface in R: **rlecuyer**, **rstream**
  - ▶ In R core: `RNGkind("L'Ecuyer-CMRG")`